

Bridgetown Research: Building AI Agents for Professional Services

[0:00] This is largely professional services, which has been my life and actually most of our team's lives for almost the last two decades.

[0:10] We are a bunch of people from consulting companies, Bain, McKinsey, and from Amazon, DeepMind, have come together to solve some professional services workflows.

[0:23] We are building agents for PE, VCs, consulting organizations and even contract.

[0:29] These are the use cases we look at.

[0:35] So overall, what I wanted to share today is essentially, obviously we haven't made as much progress as some of the other agents, like for instance coding agents that are much, much ahead in terms of benchmarks and outcomes of their delivery without human in the loop compared to where we are from a professional services agent world.

[0:56] And what has led to those?

[0:58] There are a bunch of structural barriers.

[0:59] Clearly, it's not the model, which has been seeing leaps and bounds.

[1:04] And where do we fit in?

[1:05] What do we do as Bridgetown Research?

[1:06] Because I don't think many people are aware of what we actually do.

[1:12] Moving on, just to set context, this is what I was talking about when I say agent efficacy has risen sharply across some of the That's been tremendous progress, but for professional services use cases, that's actually our own home benchmark.

[1:28] We do benchmark agent performance across 500 consultant tasks, which are specialized across like whether it's growth, whether it's intelligence, whether it's cost,

whether it's some of these IT transformation, AI transformation readiness, customer surveys, et cetera.

[1:44] These are very specialized tasks that we test or eval performance of different models on it hasn't improved as much.

[1:54] And there are two reasons for it.

[1:55] And I think we have also seen this in the benchmarks released by Harvey or Val AI, where when individually some of these tasks meet the rubrics, but they are not deliverable, which is very important in professional services world.

[2:09] These need to be output client ready that need to be sent out, that need to be decision ready so that it actually helps executives make those decisions.

[2:20] Okay, now the two barriers, I think barrier two, I won't spend a lot of time on because for any agent lab, this is a barrier we all face, which is scarcity of ML talent here across different skill sets.

[2:33] And that's also something we'll likely touch upon in terms of what we see as easing from a ML talent perspective.

[2:42] But the bigger problem that we had to deal with as an agenting lab building these professional services agents is lack of expert data.

[2:53] Actually, it looks visually off, but there's a huge bar here, which is majority of the cost when we think of building a model that on a basket of 500 consulting tasks that I showed you on the right-hand side of the previous page.

[3:08] To get to a threshold acceptance on all of those tasks, we needed to spend X amount of money.

[3:15] 67% or two-thirds of that is actually getting that data.

[3:19] And in a moment, I'll tell you what that data is and why that was so difficult for us.

[3:24] And we still continuously capture this on our ongoing projects.

[3:30] If you were to think about how this would have looked for other kinds of agents, like coding agent, for instance, where these are verifiable tasks, Compute would have been a much larger bar.

[3:39] It would have almost flipped.

[3:41] So in our space, expert data is the most expensive resource, which sort of is like some part of judgment, some part of choices made by these experts when they solve a use case end to end.

[3:57] And like I said, just touching a bit on the talent element.

[4:03] While the scarcity exists and we all know it, I think it's easing from both ends, from both demand as well as supply perspective.

[4:09] Supply, we'll have more ML scientists coming, like 15 to 20% estimates, I think roughly exist.

[4:16] But even from a demand perspective, across the stack, now training infrastructure, inference deployment, are places where talent is not as necessary for an agent stack perspective because I can outsource this.

[4:29] training infrastructure especially for like thinking machines or like fireworks through an API call.

[4:34] And inference and deployment is something that with open weights models becoming so good and being served by hyperscalers, this is not anymore that kind of constraint for us.

[4:46] Having said that, model research and training is the most important skill set that we look for and that is easing to some extent with supply increasing but from an overall talent perspective expect this to ease in the near future spending a bit of time on the bigger constraint for us which has been again this has to be built use case by use case which is for problem statement by problem statement

[5:14] for most uh consulting or investor use cases here

[5:20] Like we spent, like millions will have to be spent depending on the number of use cases you're looking to codify.

[5:27] Because you have to track not just the output, which could be a repository at any of these knowledge points.

[5:33] But this is actually what decisions were taken, what was the communication, which is not held in a repository anywhere.

[5:39] This is either over emails or on calls or like on their like notes, everywhere in folder structures.

[5:47] But what is necessary is how are they, what are the hypotheses, what are the evidence they're looking for to prove what kind of hypothesis and who are they gathering it from, what do they prefer in terms of sources, so what have they arrived at, what are the drafts, what is the partner input, for instance, like in professional services world,

[6:06] the way this would work is We would look at an investment use case, somebody spends like four weeks of their time and probably half a million building a case for whether to invest in asset or not and what are the value creation plan if you were to invest etc.

[6:22] Towards end, maybe like at the third week point somebody, a partner, associate partner would look at and say, hey, I don't think this is right, go rework it, at least some parts of it.

[6:30] And then you do the whole thing again.

[6:31] So the critique and revision, this is a very important part of the process to capture.

[6:36] So this is what we call expert trajectory, which is very important to register and then fine tune our models based on this.

[6:45] So the more trajectories a firm can actually build, a consulting firm can build this, a firm can actually build, but you need to track at every point in time the data, the process, and who is, what's the choice or preference of these individuals.

[7:01] And the way we see this easing or the way we're seeing path forward here is actually an operating platform that records everything, which means that platform needs to provide tools for strategy professionals or like investment professionals to be able to conduct their research end to end on the platform and get the critiques on the platform,

[7:21] decisions done on the platform and the final version and multiple interim versions, data, analysis, who was like

[7:28] Primary research had to be done on the platform.

[7:30] This has been our learning, which is how.

[7:32] rebuild our platform end to end to cater to these professionals.

[7:37] So that judgment gets captured live.

[7:41] The one thing however is you can only do it going forward.

[7:44] Whatever projects or decisions you take by doing research on the platform is the trajectory you can capture.

[7:51] So if you are two years in not adopting this today, that much of trajectory is lost because there is no other way that you can codify it.

[8:01] I think the last point, like everybody spoke about data flywheel, I think all of us are aware of this, which is if we can get to that kind of firm-specific future, I'm not saying we have this today.

[8:14] We are rolling out pilots where we are building firm-specific trajectories and firms fine-tuning models for specific firms.

[8:22] These are larger firms, and our model is slightly different when we work with boutique consulting firms or boutique investment firms.

[8:29] But in these cases, we see the flywheel leaning out already, and I'll show you a live demonstration towards the end.

[8:36] Fund-specific agents lead to higher quality, faster delivery, so you deliver more, you get more business as well.

[8:43] Greater adoption, more expert trajectories, which is the very important thing that's driving the data wheel, and better training and evaluation we need to do this.

[8:53] You're not actually investing upfront.

[8:54] That 67% bar is not a decision you take where you need to put in a lot of money to generate this data.

[9:00] This is now happening as you do your everyday work because that's what we are registering on the platform.

[9:05] It's saved at a very granular level so that the model can be fine-tuned based on these choices and decisions that are being made.

[9:16] Now, where do we fit?

[9:18] I wanted to show this with an example.

[9:21] For instance, a consultant is getting a problem statement.

[9:24] They have a scope.

[9:26] What are the hypotheses?

[9:27] Like, what's the, how do I even design the questionnaire?

[9:30] What are the topics?

[9:31] Oh, if for this scope, this hypothesis, this is how I want to phrase this question.

[9:35] And to whom do I want to speak to for primary research and what sources do I prefer?

[9:39] I do some secondary research and then I build an insight which flows into a storyline.

[9:45] I have a deliverable and then I get some critique.

[9:48] That would be a loop that runs and then you have a final deliverable.

[9:52] Now our platform, our platform's name is Team Room.

[9:55] That's what we deploy at these firms, whether it's corporate strategy desk or like strategy desk or even investment professionals.

[10:02] Team Room captures all of these in a stepwise fashion because there are tools that let consultants and analysts actually run the whole thing end to end on the platform itself.

[10:14] Starting with It captures on your framing or hypothesis.

[10:20] It understands how you're framing the objectives.

[10:23] It actually generates a plan.

[10:25] Lets you edit from its own knowledge repository.

[10:27] It generates the plans, these are the goals, these are the hypotheses.

[10:31] You can edit it so it understands the changes you've made.

[10:35] The chosen hypothesis.

[10:36] And it understands, okay, who are you sourcing on the platform?

[10:40] Because we also have...

[10:41] our own primary agents that actually conduct these conversations these are AI calls typically uh in a in a traditional Consulting environment or investing environment you would analysts or Consultants would do these calls themselves so to do 40 calls it would take them maybe three to four weeks of their time for a team of three to four people because you have to get on get time with these experts

[11:06] you could do that in a week week and half with AI because now experts have

[11:11] Like they take it at their convenience, they do these calls over the weekends or like on their way back home or way to work.

[11:18] You get those insights very, very quickly.

[11:19] So that saves time, but that's actually besides the point.

[11:22] I wanted to say, so who are you talking to?

[11:26] What are you covering?

[11:27] What are the most important insights?

[11:29] Where were you asking for follow-ups?

[11:31] All of that's getting captured in this, the second part.

[11:36] Then analysis.

[11:38] There are agents again on the platform that turns whatever is being received from primary and secondary, depending on the sources you choose on secondary, into like defensible hypothesis.

[11:50] It itself drafts the storyline.

[11:51] It keeps telling you this is what I'm seeing as summary insights emerging.

[11:55] End of a week, end of two weeks, end of three weeks, whatever you run.

[11:58] And we have seen people run these periodically.

[12:00] If it's like a one month refresh on competitive assessments, you can run that.

[12:05] You might think, okay, now can I do all of this on cloud?

[12:07] We do see enterprise cloud enterprise GPT being leveraged, but the thing is you can't do end to end on cloud or GPT, which is you can't conduct primary research.

[12:19] You don't have a way to steer your sourcing, which is I don't want you to get anything from this source, I want you to do it from this source.

[12:28] And you don't have all the documents in a repository.

[12:32] the agent having access to it.

[12:34] You'll have to upload everything, you have to choose and upload.

[12:38] And storylining, this is a purpose-built harness for strategy professionals.

[12:41] So the storylining is the way you like to see it.

[12:44] And it also creates draft pages against the storyline with an assertion backed by evidence and you can trace it.

[12:51] Everything is traceable on the platform.

[12:52] If you have like an assertion or an evidence, you can click on it, it takes you to a transcript or a document.

[12:57] from which it's gathering the insight.

[12:59] So this whole thing as a harness is purpose built.

[13:05] So essentially the whole choice that it's gathering and the output, that's what compounds the learning here, feeds into, like that's how we fine tune our models.

[13:18] And the compound is private as a single tenant model for large companies.

[13:24] And I think I've covered to this to some extent covered to some extent but this is primarily what underpins our platform there are primary research agents secondary research agents and analysis and reporting which enables anybody to solve a business problem end to end you can conduct these conversations you can do live calls also on the platform not just AI conversations but we have seen 90% of

[13:47] our conversations uh being AI and like 5 to 10% are live calls where you're just trying to understand unknown unknowns because you can't do that with

[13:56] an AI agent.

[13:57] So known unknowns is something you can poke further when you do primary calls.

[14:02] It's extremely, it makes a consulting job extremely, I wouldn't say simple, there's a ton of judgment required, a human in the loop is very, very important.

[14:11] But it enables somebody to focus their time on the most important things and not waste time.

[14:17] For instance, it generates a question guide already for you.

[14:19] If you just say this is my objective, it identifies goals, generates a question guide, It showcases the kind of experts you can talk to.

[14:26] You can just say, hey, yeah, recruit this person for me.

[14:29] It gets the person to interview.

[14:31] It interviews the person, gets a response back.

[14:35] You see the transcripts in the repository within your project context.

[14:39] Secondly, we have a bunch of sources.

[14:41] We have a repository of data and transcripts.

[14:45] We do also, we have the data as well.

[14:49] So we have our partners that run these projects, live projects, and share all of that information with us.

[14:56] So we have a bunch of data on industries.

[15:00] that we focus on, which is we are expanding into more industries today.

[15:04] We look at technology, healthcare, industrials to some extent.

[15:07] And analysis reporting, here's again, it's very, I wouldn't say it doesn't create 100% of the output, which is where the self-learning loop needs to kick in.

[15:17] Or we need to further fine-tune for some of these cases.

[15:20] Some of these cases, it's pretty good.

[15:23] I would say it's almost like 70 to 80%.

[15:26] answers that you can turn into outputs are provided from the platform, but not for all use cases at the moment.

[15:36] Yeah, here's some, like I said, some workflows or some agents we have built for our customers across for investors.

[15:46] Thesis development, origination screening, P&L prep.

[15:49] This is where I actually see a lot of usage on the platform.

[15:53] Post-LOI due diligence for mid-market funds that typically don't go to like NBVs for their due diligences, leverage our platform.

[16:04] Most investors I see leveraging the platform, like large investors leveraging the platform for pre-LOI prep because that's something they do themselves or so far have not had credible alternatives to do this for them.

[16:17] Now what this is doing is it's expanding the top of the funnel in terms of how many signals, how many assets you can get like preliminary signals on.

[16:26] That's step three, I see a lot of traction on the platform.

[16:30] Due intelligences, again specific sets like mid-market and lower mid-market funds, use that quite a bit.

[16:35] Value creation and transformation is where you're seeing increasing awareness and push because work with a lot of investors and investor value creation are the use cases we do a lot of pilots in.

[16:47] Whether it's a strategy, AI maturity diagnostic, again, I think most people spoke about this.

[16:52] This is something that's seen a lot of traction as well.

[17:00] Yeah, a quick case study, I would say, of this is a...

[17:05] mid-sized tech consulting firm based in the US.

[17:08] The majority of our go-to-market is North America and Europe.

[17:11] We do very little in India.

[17:14] Obviously, we work with some of the big four, so we do global for them.

[17:18] But otherwise, our focus has not been on Indian market as yet, but it doesn't change a lot.

[17:23] The use case doesn't change a lot.

[17:27] This company, we've deployed our platform over the last almost a year, 200 live projects delivered.

[17:35] 3,000 plus transcripts generated as part of it.

[17:37] Transcripts is, although not the core thing, it's actually about delivering end to end with a human in the loop.

[17:45] And the interesting thing I wanted to highlight is the impact that it's had.

[17:50] We have seen team sizes move from an M plus four to even M doing-- sorry, M, manager plus four people in the team to half a manager and maybe two people.

[18:03] two analysts, or sometimes even like one analyst doing this.

[18:07] And it's a commercial due diligence which would take like a manager and three to four people to work on it for like three weeks.

[18:15] It's two weeks faster.

[18:16] It's been done in like, I've seen some projects end in a week on the platform.

[18:21] It's one and a half to two weeks on an average.

[18:23] And it costs a third.

[18:27] This is again evolving a lot of Judgment based use cases or like enterprise transformations have some way to go because you have to capture a lot of judgment and they'll also have much more the human layer or interference will have to be higher to start with.

[18:42] But then as the model learns and understands the processes, who's doing what, it should be able to pick that up.

[18:50] And these are like some of the, we also build a lot of skills for our customers where I think in this case, for the customer, because They have started with us for their CDD, for their diligence practice, which is the largest for them, based in North America.

[19:08] So here is what we are seeing.

[19:11] The world, I think, will always need judgment, consulting judgment.

[19:17] That's not going to go away.

[19:18] But it's getting a lot more efficient.

[19:21] And when I speak to-- because we also cater to consulting firms, we set the platform to consulting firms.

[19:29] Large firms, for instance, it's their model and we don't have that data.

[19:34] But it's also opening up, I think like we've heard in the previous session as well from Sukhi, it's also opening up more opportunities for them, which did not make economical sense before to hire consultants.

[19:47] And consultants could not cater to those kind of opportunities.

[19:49] For instance, pre-LOI, someone can actually deliver it in like two, three days.

[19:54] leveraging the platform and getting charged a fourth of what they were charging, which makes economical sense for their end customer.

[20:01] So we're seeing some of these, we're seeing some of new innovative workflows like tech diligences, which we were not doing before, but because we can hook into code databases and analyze, we have scripts to like our workflows actually enable and deliver operating like, tell me about metrics like what's their code deficit like, what's their, you know, time to deployment and stuff like that.

[20:27] So there are standard tech TDs that are being done on the platform.

[20:31] So this was something that consulting companies were not doing before.

[20:34] So it's also expanding their aperture while improving their productivity.

[20:39] But having said that, maybe some of the traditional billing rates will be under pressure for sure.

[20:44] That's what we are seeing.

[20:49] Yeah, outside diligences, we've also had like processes being laid out with the kind of primary agents we have.

[20:57] This was a case study where somebody was able to do more than 1000 interviews in a couple of weeks to actually scope out, I think this was Salesforce implementation if I'm wrong, but to scope out a system of record, a system of engagement implementation at an end customer.

[21:14] along with human group on top of it, of course.

[21:16] But they would not have been able to conduct the kind of interviews they have conducted across geographies in a couple of weeks.

[21:26] Yeah, that's that from us.

[21:36] It was fascinating.

[21:37] I wish this was available 10 years ago when I was spending 16-hour days on VDDs.

[21:44] So, few questions from the audience.

[21:46] I think the first and probably most repeated one is, how do you manage the quality?

[21:51] Because unlike coding agents, which have specific ways to test, especially if it's not specific quality, how do you ensure it's there?

[22:01] So, we do have rubrics for each of these consulting tasks that we evaluate these on, not just how comprehensively something is answered, what's the relevance, what's the factual accuracy, what's the logical consistency.

[22:15] And we have humans to actually evaluate this.

[22:18] And that's how we get to those benchmarks.

[22:20] And we also have, we provide options for the end customer in those firms, and they choose one or the other.

[22:26] And that's how our model learns, and that's how we evaluate as well.

[22:29] We conduct, we design these evals, we have some rubrics in place to evaluate the quality.

[22:34] Because it needs to be verified.

[22:36] So you need a human to verify the quality of the output.

[22:40] Right.

[22:40] So-- For the 500 consulting tasks that I was talking about.

[22:43] I understand.

[22:44] So but at every-- let's say, every client project used to have a human verifying almost each piece of delivery block.

[22:52] Yes, because there is a human in the loop.

[22:55] Even if I-- let's say, PwC is using this platform for their delivery, they're going to

use the whole A consultant is going to kick off all of this research.

[23:04] He's not going to share the final output of the PPT that the agent is producing directly.

[23:10] So there is going to be human in the loop there.

[23:13] But all the corrections are being done on the platform.

[23:16] So the platform is also capturing the choices and preferences of people there.

[23:21] That's how the learning happens.

[23:22] Similarly, even at the question guide level, it takes the first stab at generating a guide.

[23:27] All the changes are being noted, registered.

[23:30] storyline, same thing.

[23:31] It generates a storyline, but whatever edits you make, so over time, the edits go down for a use case.

[23:39] I guess that learning is also firm-specific, so you kind of know what a firm requires.

[23:45] Yes.

[23:46] It's not at a test level, but we're doing it at a firm level.

[23:49] At a firm level, OK.

[23:50] So then, related to it, I guess, is how much of the performance gain comes from capturing the initial research process versus the more critique and revisions that follow.

[24:02] Where is the gains?

[24:03] Yeah, very interesting question.

[24:05] So when we deploy something like a base model, you see about 40% that can go out as an output.

[24:13] Like a PPT can be accepted 40% of the time, but the rest of the time, somebody is doing something on it.

[24:17] Same is the case with storylines, same is the case with.

[24:22] And then once the-- because that's where our tribal knowledge is fed in.

[24:27] We also have a group of people, our partners, that are delivering these on the platform.

[24:32] And hence, we've already learned it before we deployed the platform.

[24:35] Post-deployment, we see at least 20%, 30% more improvement as it's capturing information live.

[24:43] Oh, wow.

[24:44] I would say like 35%, 40% of the.

[24:45] Oh, nice.

[24:49] Then, I think there is just one more question on-- actually, interesting.

[24:55] How does this change your talent development?

[24:59] Because so much of it is of productivity is gained.

[25:04] How do you get the critique?

[25:08] I think about it too, right?

[25:09] Because it's the reality, which is you don't need going forward.

[25:17] You don't need the kind of pyramid consulting companies have been managing historically.

[25:24] But having said that, you do need the talent to be able to become partners, associate partners.

[25:28] So I think the pyramids are going to be narrower, but it's not going to be like a one single end.

[25:34] But you don't need like 4 to 5 people for every manager in an organization.

[25:38] You probably go to like 2 to 3.

[25:40] That's how I think about it, because you genuinely don't need so many analysts or associates, because all of those activities can be done.

[25:49] A lot of non-important, like decision-making time is like, It's solved, and it will be solved very soon.

[25:57] OK, but that's a very different model versus the traditional consulting model?

[26:02] But you still need people to gain expertise, right?

[26:04] So they will have to go up the ladder.

[26:07] They are not just going to burn out of the blue.

[26:10] But you will not need-- I almost visualize it as like, maybe it's going to be like a cylinder and then like a thin thing on top, because you don't need the whole fat base.

[26:21] Interesting, interesting.

[26:22] I, those were the main questions.

[26:24] Thank you very much.

[26:25] I'll do it.

[26:25] Thank you.

Notes powered by [CraftNote](#)