

Session Summary

Overview

- Date noted for this summary: **September 4th, 2026**.
- The session featured **Prashant** (Co-founder and CEO of **Cure.ai**) discussing the company's evolution, product strategy, clinical impact, regulatory considerations, and commercialization model.
- The core theme was that AI in healthcare is most valuable when it is embedded into real clinical workflows to improve **early detection**, **follow-up management**, and **guideline-concordant care** rather than merely generating risk scores or standalone predictions.

Main Topics Discussed

1. Cure.ai's Mission and Global Reach

- Cure.ai's stated mission is to help clinicians detect disease earlier and close major gaps in care delivery.
- The company has built healthcare AI solutions across imaging, text, and voice.
- Key scale metrics shared by Prashant:
 - **1.5 billion images** used for training
 - **26 FDA-cleared algorithms**
 - **65+ CE-certified indications** and additional global validations
 - **250+ papers** published
 - **40+ patents** mentioned in one part of the talk, and **35+ patents** in the intro
 - **150+ engineers**
 - **45 million+ lives impacted**
 - **5,200+ hospitals/sites** using the platform across **107 countries**
- The company's solutions are deployed in both developed and resource-constrained settings, including **Africa, Latin America, Southeast Asia**, as well as the **US and Europe**.

2. Why Early Detection Matters in Healthcare

- Prashant emphasized that many diseases leave signals long before symptoms appear, especially in medical images such as **chest X-rays, CT scans, and MRIs**.
- He highlighted lung cancer as a major example:
 - Late-stage lung cancer has about **5% 5-year survival / 95% mortality**.
 - Early-stage lung cancer can have about **80% 5-year survival**.
- The company's view is that the biggest opportunity is not just "cure cancer with AI," but rather improving real-world healthcare operations so that patients are diagnosed and treated earlier.

3. The Two Major Failure Points: Missed Findings and Missed Follow-Up

- Prashant described two major places where patients fall through the cracks:
 1. **Image-level misses**: findings missed on the scan itself.
 2. **Follow-up-level misses**: findings that are reported but never acted on.

- He stated that for scans like X-rays, CTs, and MRIs, there is roughly a **5%–10% chance** that an important finding is missed.
- He also stated that approximately **60% of incidental findings are not actioned on**.
- A major example was **lung nodules**:
 - About **1.6 million** patients have an incidental pulmonary nodule documented in a radiology report.
 - Only about **600,000** receive guideline-supported follow-up.
 - Roughly **1 million** patients fall through the gaps.
- He said there may be **10–20 million** patients across incidental findings in the US who could benefit from automated detection and follow-up support.

4. Clinical Variability and the Role of AI as a Second Reader

- Prashant discussed inter-reader variability in radiology:
 - Different radiologists can disagree on the same chest X-ray or chest CT about **35%–40%** of the time.
 - He described AI as a “second reader” that can increase accuracy, citing an improvement from **68% to 82%** in one example.
- The goal is not just detection, but also reducing inconsistency and improving reporting quality.

5. Cure.ai Product Stack: Q-Lens, Q-Track, and QRX

- Prashant outlined Cure.ai’s product layers:
 - **Q-Lens**: image-based AI that detects findings on scans, such as nodules, small bleeds, TB signs, or pneumonia.
 - **Q-Track**: workflow and follow-up management layer that ensures patients with suspicious findings actually receive appropriate next steps, such as referrals, CT scans, biopsies, or pulmonology consults.
 - **QRX**: treatment and guideline-concordance layer that nudges clinicians toward the most appropriate therapies based on multimodal data and guidelines.
- He stressed that a risk score alone is insufficient unless the system also drives actions, referrals, documentation, and reimbursement-ready workflows.

6. Multimodal Healthcare AI and Workflow Integration

- Cure.ai now integrates multiple data types:
 - **Radiology images**
 - **EMR / EHR data**
 - **PACS data**
 - **Clinical notes**
 - **Scribing / doctor-patient conversation data**
- The company builds “referral-ready packets” that assemble relevant evidence from scans, reports, prior imaging, labs, pathology, and notes.
- Example given: a lung nodule that grew from **6 mm to 9 mm over a year** could be surfaced through combined imaging and text signals, even if no single data point alone triggered action.
- The emphasis was on ensuring the output is immediately usable in systems like **Epic** and integrated into hospital operations.

7. Guideline-Concordant Care and Reimbursement

- A major point throughout the talk was that in the US healthcare system, guideline adherence is closely tied to reimbursement.
- Cure.ai's tools help ensure patients receive guideline-concordant next steps such as referrals, scans, or therapies.
- Examples included:
 - A lung nodule patient receiving a pulmonology referral and CT scan with correct paperwork for reimbursement.
 - COPD patients who qualify for biologic therapies but often do not receive them; Cure.ai nudges physicians at point of care to improve adherence to clinical guidelines.
- Prashant said that only about **15 out of 100 eligible COPD patients** receive certain biologic therapies, illustrating the care gap.

8. AI Agents and Operational Follow-Up Automation

- He described an agentic workflow in which an AI system inside the hospital can:
 - Review referral packets
 - Reach out to the ordering physician / primary care provider
 - Contact the patient via voice agent or nurse navigator
 - Schedule specialist appointments such as pulmonology visits
- The objective is to close execution gaps, especially in incidental findings, where no one may be actively managing next steps.

9. TB Screening and Low-Resource Settings

- Cure.ai's origin story centered on **tuberculosis (TB)** screening in low-resource settings.
- TB statistics cited:
 - Around **11 million** cases annually
 - About **1.5 million deaths** annually
- The first customer was a **mobile van-based screening program in the Philippines**.
 - X-ray reads that previously took **4–5 weeks** were reduced to about **30 seconds** using AI.
- The **WHO** later issued guidance in **2021** supporting autonomous AI use for TB detection on chest X-rays.
- The company has deployed across **100+ countries** and **3,000–4,000+ hospitals/sites** for TB-related screening.
- He also mentioned deployments in difficult environments such as boats, remote areas, and **Mount Everest Base Camp**.
- The goal in TB has been to reduce diagnosis time dramatically, with one example cited of bringing diagnosis time down from **2 weeks to 2 hours**.

10. Extending Beyond X-ray: Ultrasound and Voice Tools

- Cure.ai is also working on AI for **point-of-care ultrasound** for TB detection.
- Gates Foundation support was mentioned as part of this expansion.
- The company developed **Aira**, a mobile/tablet-based note-taking and decision-support tool that listens to healthcare worker–patient conversations.
- Aira's functions include:
 - Converting conversations into structured data
 - Nudging healthcare workers with next questions or next steps

- Supporting treatment and investigation recommendations
- Prashant also framed Aira as a **pandemic preparedness** and surveillance tool because it helps understand disease patterns from frontline encounters.

11. Data Flywheel and Model Improvement

- Cure.ai processes roughly **15–20 million patients annually** through its systems.
- The company collects not only scans and records, but also feedback on whether clinicians accept or reject recommendations.
- This creates a feedback loop where every click and action helps improve future models.
- Prashant described this as a flywheel effect: the AI improves itself through real-world usage and clinician behavior.

12. Regulatory, Liability, and Privacy Considerations

- **Liability:** Prashant said that the liability generally remains with the physician or healthcare organization making the decision, not the AI model itself.
- He noted that AI is supporting decision-making rather than replacing clinical responsibility.
- **PHI / privacy:**
 - PHI is typically stripped when possible for decision-making.
 - Patient identifiers are needed only for linking results back to EMR or PACS records.
 - Cure.ai operates with local data infrastructure to comply with national laws.
- Examples of geographic compliance infrastructure mentioned:
 - **US AWS cluster** for the US
 - **India AWS cluster** for India
 - Use of other local cloud providers where necessary, such as **Alibaba** in some regions and **VSN** in Vietnam
- He referenced compliance with **HIPAA**, **GDPR**, and local national laws, with frequent audits.
- On regulatory risk, he said the risk is lower when the company works with physicians and provides decision support rather than direct-to-patient autonomous decisions.

13. Commercial Model, Pricing, and Regional Strategy

- Prashant stated that deployment prices in **India** are about **10–20x lower** than in the US.
- Cure.ai's primary go-to-market focus is the **US**, while India is mainly a TB market for the company.
- He said the value of these workflow-heavy products is higher in the US because of broader **EMR adoption** and more structured reimbursement pathways.
- In India, the value outside TB is lower because the infrastructure and workflow maturity are different.

14. Economics of Early Diagnosis for Hospitals and Pharma

- Prashant explained that early diagnosis creates value for both hospitals and pharmaceutical companies.
- Example economics shared:
 - A single lung cancer patient may be worth about **\$100K** to a hospital in downstream services.
 - The same patient may be worth **\$200K+** to a pharma company, potentially **\$200K–\$400K** over 1–2 years of drug therapy.
 - One month of lung cancer drugs may cost about **\$17,500**.

- This creates incentives for both healthcare systems and life sciences companies to invest in earlier detection and follow-up.
- Cure.ai works with nearly all of the **top 10 life sciences companies** globally in some capacity.

15. Partnerships and Distribution

- Two major distribution partnerships were highlighted:
 - **Microsoft**: helps distribute Cure.ai products.
 - **Intuitive Surgical**: an exclusive distributor for some products, especially tied to pulmonology workflows and the **Ion** bronchoscopy/biopsy ecosystem.
- Prashant also referenced a partnership with **Sigla** (or similar, based on transcript quality) for early COPD detection in Indian hospital systems.
- In these arrangements, partners generate leads and build relationships, while Cure.ai handles implementation, integration, sales completion, and deployment.

16. Product Design and Ease of Adoption

- Prashant emphasized that clinicians are busy, so the software must not add work.
- Cure.ai's product is designed so that:
 - Pop-ups appear only when relevant.
 - Clinicians can click to open details if needed.
 - A full DICOM viewer is available for review.
- He said training is generally simple and can be done in about **one day**.
- The company tries to minimize workflow burden to improve adoption.

Key Takeaways

- Cure.ai's strategy is centered on moving beyond simple diagnostic AI toward a full clinical workflow platform that identifies disease, triggers action, supports guideline adherence, and helps with reimbursement-ready care delivery.
- The strongest value proposition is in settings where there is enough digital infrastructure to support integration, especially in the US.
- The company's most important differentiator appears to be its combination of **deep clinical AI, workflow integration, multimodal data, and regulatory breadth**.
- TB screening remains an important global-health use case, but the company's commercial and workflow-heavy growth is increasingly tied to the US and pharma partnerships.

Follow-Up Points

- No explicit follow-up meeting was scheduled in the transcript.
- The closing segment suggests ongoing discussion around future expansion into:
 - **Molecular diagnostics / biomarkers**
 - Additional multimodal workflows
 - Continued expansion of guideline-concordant therapy support
- The host indicated there were more questions, but the session ended due to time.

Action Items

- Continue scaling hospital deployments in the US, with a focus on restoring growth toward the stated target of **700 sites by the end of the year (December)**.
- Expand multimodal integration across imaging, EMR, PACS, voice, and clinical notes to improve follow-up and treatment automation.
- Further develop **Q-Track** execution workflows and **QRX** treatment-guidance workflows.
- Continue working with pharma partners on early diagnosis, treatment protocol optimization, and clinical trial support.
- Expand point-of-care and low-resource use cases, especially for TB and ultrasound-based screening.
- Maintain ongoing compliance and auditing for HIPAA, GDPR, and local-country data privacy laws.
- Improve clinician-facing UX to preserve low-friction adoption and keep training time minimal.

Action Items

- ☐ Share the integration plan for the US hospital deployment workflow — @Prashant Bhai
 - ☐ Provide the details of the Intuitive Surgical distribution partnership — @Prashant Bhai
 - ☐ Prepare the follow-up answer on expanding into molecular diagnostics and biomarkers — @Prashant Bhai
-

Transcript

Host (0:00)

AR technologies with solutions that offer early detection and management of lung cancer, neurocritical findings, and infectious diseases to support clinicians and proper developments in pharmacy and medical device industries. Under his leadership, Cure.ai's specialized healthcare solutions has 26 FDA indications and over 65 CE-certified indications and other global validations, reaching more than 5,000 global healthcare and multimillion patients. This includes remote resource-constrained primary care settings in developing countries such as Africa, Latin America, Southeast Asia, and Western nations such as US and Europe facing clinical burnouts, backlog, or bottlenecks. Prashant is a global speaker and a recognized thought leader at the intersection of artificial intelligence, healthcare innovation, and global health equity. In 2025, he was named a Bloomberg Economic Catalyst, recognizing global innovators shaping a more resilient and inclusive future, along with being named one of Forbes India's Top 30 AI Minds, an Economic Times India Top Innovator, and double recognition as a Time 100 Most Influential Company for using AI to reshape access and diagnosis of high-reliant diseases. Prashant holds more than 35 patents and earned his PhD and MS in Operations Research from Georgia Institute. Along with a bachelor's degree from IIT Delhi, we are very happy to host you, Prashant, today for the conversation. And please, stage is all yours.

Prashant Bhai (1:46)

Thank you so much. Hi, I will thanks for inviting me. Thanks for the great introduction. It was actually way more than I expected. I just my name is Prashant Bhai. I am co-founder and CEO of Cure. And I'll talk about the cure journey and what we have built, talk about healthcare in general, but I'll start with a video.

Speaker 3 (2:31)

We cannot see until it's too late. Diseases that affect millions every year, often detected too late. Because the truth is, many diseases leave signals long before symptoms appear, hidden inside medical images, chest X-rays, CT scans, signals we often miss, but signals that AI can see.

Prashant Bhai (2:49)

Once was a cold that had for a little six months. Not yet, but it became almost not GPU friendly for an X-ray.

Host (2:56)

So it's sinambil mo pa X-ray do kami ni nalamam ko na na pinatawag kami dalawa, pinatawag kami yun o sinabihan.

Prashant Bhai (3:06)

And that's why screening matters. If you're able to diagnose early, you could have extremely high survival. And that's that's sort of the problem we're trying to solve.

Speaker 3 (3:15)

And that's what Cure AI is built to do: solve one of healthcare's biggest challenges—detect disease earlier.

Prashant Bhai (3:21)

Biliyon tigil di ako na diseases, di ako na ailments, di ako na chest. Problems that we can leverage on these, and we substitute with with some updates in the artificial intelligence to be able to profile other diseases.

Speaker 3 (3:38)

Cure AI is driving deployment across geographies through local and global partnerships.

Prashant Bhai (3:42)

Proud to be with this association of Cure AI and SLA. We are going to recognize these cases very early and identify patients who have a blockage in the in one of the main arteries in Stockholm, and within 24 hours we are able to recognize those patients using this technology.

Speaker 3 (4:02)

From the Mahakam to the highest point of Everest, from Nigeria to the United States, from screening vans to stroke emergency rooms, AI is helping clinicians see earlier and act faster across hospitals, across health systems, across the world. Cure AI deployed across 107 countries Powering thousands of healthcare sites and reaching millions of lives.

Prashant Bhai (4:25)

For you, we are working towards one goal: towards positively impacting more than a billion lives through the power of care. Yeah, so I, uh, wanted to quickly talk about AI in healthcare in general, right? I think there's a lot of hype around AI and in healthcare specifically that talks about solving cancer, curing cancer completely. But actually, I mean, if you look at the real challenges, I mean, we have been reading cure for the last 10 years. The real challenges that we observed are actually integrating AI into the healthcare workflows, making those workflows better with AI, and I'll talk about some of those where the gaps are. And that is where a tremendous value can be created. If you diagnose a lung cancer patient 1 year earlier, 2 years earlier, they could survive 20 years longer, 25 years longer. So survival rates for late-stage lung cancer is 5%, 5-year mortality is 95% is the mortality for late-stage lung cancer. For early-stage lung cancer, that is about 80%. I mean, 80% of people will survive after 5 years if you diagnose stage 1, stage 2. And that is the kind of problem that we're trying to solve. We started with lung cancer, TB, but now we go beyond that and we're solving problems across multiple disease areas. And talk about the real issues, right? So if you look at it, I mean, there are two places that you are missing patients, right? I mean, when you go and get a scan, get an X-ray or CT scan or an MRI, there is a five to ten percent chance that important finding, not of any finding, important finding is missed on that scan. And that is where AI comes in. AI can find that that missed finding and can help improve the quality of the radiology reports, right? So this is missed on the image. Second is actually even bigger, right? The fact is that you report, I mean, you report that there is a nodule on a chest x-ray, and I might be going a little bit more technical on the clinical side, and I mean, I try to be more sort of try to help you understand better, but the nodule is something which can lead to a lung cancer, right? And if you miss a nodule, it could be lung cancer in 2-3 years, and that is what we're trying to solve. But another problem is that somebody reports a nodule, radiologist reports that there is a nodule on a chest X-ray or chest CT, but there is no action on that because that— and especially I'm talking about the US health system— these are there in the reports, nobody's mining the reports, nobody's creating a follow-up, nobody's saying that this patient needs to get a CT scan or a biopsy, and that patient does not know that we need to take further action. And that is a huge number, right? 60% of incidental findings are not actioned on. So they don't get any follow-up, and that's again a huge problem that we are trying to solve. So misses at the image level and misses at the follow-up level. And these are certain numbers, right? I mean, so if you look at it, 31%— if you take a simple chest X-ray, I mean, all of us have taken chest X-rays, you get another radiologist

doing the chest X-ray and you have a report. You get another radiologist to read that chest X-ray, they will disagree with the original report almost 35-40% of the time. Same is true for a chest CT, same is true for most findings. So imagine, I mean, 2 radiologists reading the same scan, they will only agree 65% of the time or lower, right? And that's a, that's a big deal, right? And that's where the AI can play a big, big effect, is that the second grader—it can be a second grader, and it ensures that you are increasing accuracy rate from 68% to 82%. being that second reader on that scan, an X-ray or a CT scan. And second is not, not just—I mean, again, I'll talk about the second leak, right, which is where it has been reported but nobody's following up on that. And that is where again AI can play a big role in sort of ensuring that you can follow up that patient. A patient who has a nodule on a chest X-ray, it is lying in the EPIC report. It is lying in radiology report and that needs to be pulled. Somebody needs to ensure that, that patient gets a follow-up referral, pulmonology referral, that needs to be reimbursed by the insurance. And creating a guideline-concordant referral is a challenging situation, and that is what we enable. Yeah, and I, I just talk about one, one specific use case, right? This is the lung nodule, right? I mean, there are 1.6 million patients with an incidental permanent nodule documented in a radiology report, right? And only about 600K receive a guideline-supported follow-up. So this is 1.6, the number of patients who should get a follow-up. So 600K is the number of people who actually got a follow-up, got an endocrine follow-up, and a million patients fell through the gaps. They did not get that follow-up. And I'll talk about the numbers, right? I mean, basically, I mean, these are at least somewhere in the range of 10 to 20 million patients across all incidental findings in the US that can be followed up, can be diagnosed. Every patient, I mean, again, it's not just the patient, it's also value for the provider system. That has been created, and this is large numbers. I mean, this is just one million is the number of lung nodular patients. Ten million will be all patients across all findings. Diagnostic delay of at least twelve months because they're not followed up on time. And if you use a technology that closes the gap, it could be ninety-nine percent plus incidental findings that are found on a scan. Can get resolved. And this is where Cure comes in. We are today in about—trained on about 1.5 billion images, uh, 26 FDA cleared algorithms, uh, 250+ papers. I mean, again, highly science-driven organization. 250+ papers that we published, 40+ patents, uh, 150+ engineers, 45 million+ lives impacted, and 5,200+ hospitals that are using us across these 107 countries. And this is—I, I'll quickly talk about this. This is the problem that we're trying to solve, right? There are 3, 3 levels that, as I just spoke about. One is that things get missed on the scan, on an X-ray, CT scan, MRI, and that is where our Q-Lens product comes in. That is finding things on a scan which humans would not have found. For example, finding a nodule. Finding a small bleed in a brain CT scan, finding signs of tuberculosis or signs of pneumonia. And that's basically vision AI. But what we also did, I mean, again, we were a radiology AI company 3-4 years back, but we realized that there are multimodal segments now and we can also incorporate language data. So when we work with health systems, we integrate with the medical records, so EMR, we integrate with the PACS, which is all the images. We integrated the scribing system. So like when you have a Microsoft Copilot or some of those, they have—they're taking doctor-patient notes or doctor-patient conversation. We combine all of that and through all of these multimodal signals, we are able to say that this patient has a lung cancer risk, for example, right? That's one example of things that we do. But that itself is not enough, and that is what a lot of people are doing, is to create that risk score saying that this patient has a risk. But telling a hospital system that this patient has a risk is not enough because they have too many things to do, and we have to ensure that there is a guideline-concordant follow-up that happens right now. Think about it this way: that if I mean I get if one patient gets a risk score saying that this patient has lung cancer, they cannot just go and get a CT. That will not get reimbursed, right? So you have to create the right set of reference, right set of reimbursement documents so that that that patient can actually go and get a follow-up. A follow-up for a lung cancer patient typically is a pulmonologist consult, uh, and a CT scan, right? So, or a pulmonologist consult and a biopsy. So what we do is we integrate into the Epic systems and ensure that those are all ordered, right? The CT is ordered, the pulmonologist is ordered, and also the paperwork for that reimbursement is created so that any patient who gets that, the hospital is—I mean, the patient actually gets reimbursed, the hospital gets reimbursed for that, and that care is delivered. So So a lot of times what happens is you set up a referral appointment, that appointment does not happen because the insurance company said this patient is not eligible for this particular treatment or this particular investigation. So that's what Q-Track does. It basically solves for execution gaps. It ensures that the patients who are at risk are being followed up, we are placing orders for a CT scan, for a referral. And manages the primary care provider communication as well. And lastly, this is again something that we are starting to build out now, right? Which is that once you find that this patient is diagnosed with a certain condition, for example, I'll give you an example of COPD, right? COPD has this latest drug which is around biologics, biologic drug, which is multiple companies have biologic drugs. For COPD, and it is very people who are eligible for that are people who

have tried triple combination therapy and have not succeeded in that. Right now, there are if there are hundred patients who are eligible for that particular therapy, only about fifteen receive that therapy because again guideline-conformant care is not being delivered. So guidelines, if you enforce guidelines, ensure that you show the doctor that this patient at the point of care, if you show the doctor that this patient is eligible for this biology. they can actually prescribe more biologics. And that's again something we are doing to ensure that with all the data that we have from imaging and text, medical records, and the notes from doctor-patient conversations, we can actually add— when the doctor is having that conversation with the patients, we can suggest that this patient would benefit from this particular therapy for COPD, colon cancer, for all kinds of oncology cases as well. And that actually increases the number of people who receive guideline-concordant therapy. So, And that is what we call the QRX product, right? So if you look at it, I mean, the top block part of it, right? 100 scans coming, 50 are follow-up indicated, 20 are follow-up completed, 10 are guideline-concordant care, right? We are increasing that. That 100 needs to go to at least 50, and we are sort of solving for each gap in where the care is delivered or care has not been delivered to the guys. And just, uh, this is fairly complex, but how do we do it, right? There are so many signals in the hospital, right? There is radiology, prior imaging, pulmonary function test, labs, pathology, clinical notes, right? We combine all of that into creating one referral-ready packet, and then this is this kind of summary that is created, and based on this, the patient gets a referral, right? This is for a lung cancer patient whose nodule grew from 6 mm to 9 mm over multiple— over a year, for example. This data was not missing. It was all just coming together, right? And that is the value of AI. You're able to bring multimodal data together and infer, uh, make inferences from that. And, uh, just to give you a sense, I mean, we do a lot of different findings, so Across neurology, respiratory, oncology, musculoskeletal, cardiovascular, hepatology. All of these are FDA cleared findings that we can detect on a scan. And you can also not only we have to detect on a scan, you can also detect on the radiology report that has been created. So like I said earlier, it's not just about finding something on a scan which a radiologist missed. A lot of times a radiologist might have found something, but that is not followed up on. And this ensures that you're following up on those patients as well. These are the capabilities that we can detect and ensure that a guideline-concordant follow-up happens for that patient. And this is how we ensure— I mean, so think about it this way, that there's an agent that is looking at all that data, and agent sitting within the hospital system, uh, it takes that referral packet from the previous page, and we saw that referral created. It takes that referral packet reaches out to the ordering physician, the referring primary care provider, to the patient, calls the patient. Again, the call can be done by an agent, voice agent, or call can be done by a nurse navigator. And lastly, reaches out to the pulmonology department, sets up a pulmonology appointment. Um, this, this is where a lot of gap is being solved because a lot of this does not happen, especially for patients who are diagnosed incidentally There is no real follow-up action in the hospital. And lastly, on guidelines on burden care, this is the example of COPD that I was speaking about earlier, right? There are guidelines for when people, when patients should be given a certain therapy. But those guidelines are not necessarily followed across across health systems, across doctors. And if you can nudge the doctor at the point of care that these are the guidelines, this is the therapy that will best work for this patient, it can substantially improve guideline concordance. And again, I talk about guideline concordance because it is in the U.S. health system context, it is super important. If you don't follow guidelines, you're not getting reimbursed. So everything. There are guidelines from the radiology, ACR, the guidelines from American Heart Association. Everybody has their own guidelines for what therapy, what investigations are recommended. And so as long as you can ingest those guidelines and the agents that we have, AI agents that we have, ingest those guidelines, ingest patient data, and ensure guideline-concordant follow-ups and guideline-concordant therapy, it makes sure that whatever is being prescribed to the patient is then reimbursed by insurance companies. And this is the example of a guideline-compliant match for a COPD patient, and that is how it will look like within a Copilot or within an Embridge kind of voice, within any kind of conversation system that hospital is using. And this is briefly where we are. I mean, we sort of started from zero in the health— in US health systems Last year we had more than 700 sites. We have more than 300 sites now. We'll be 700 by end of this year, by December this year, and growing, growing very rapidly across all of this portfolio. 2 things that help us again from a distribution perspective: we have a partnership with Microsoft, they distribute our products, and then we have a partnership with Intuitive, which is the surgical system, automated surgery system. And they also distribute our solutions across health systems in the US. They are an exclusive distributor for our products. And this is sort of where we are. I mean, we are in some of the top health systems— UPMC, Mount Sinai, Orlando Health, WellStar. WellStar is the largest lung cancer program in the US. So we work with a large number of systems across the US and 2 million families back in India. Now when you think about what we are speaking here, right, I mean, finally when

you diagnose a patient early, uh, there is a lot of value that is created for pharma companies as well, and they are also quite incentivized in improving diagnosis. If you look at it, I mean, a lot of pharma companies today, life sciences companies, are investing into early diagnosis, and that is something that we started working on about a couple of years back where we work with pharma companies around optimizing treatment protocols, that end-to-end purchase that I showed earlier. We work with them around early lung cancer diagnosis, early COPD diagnosis, early TB diagnosis. So they sort of subsidize deployments in health systems, and we also work with them in optimizing clinical trials, so patient selection for clinical trials, data, real-world evidence, multiple things that we work with pharma on. And all of these, uh, are, um, use cases that are extremely valuable for pharma companies. And today, that we list at the top, uh, are companies that we work with— I mean, again, almost all of the top 10 life sciences companies in the world. And lastly, this is where we started. I mean, so our, our US, uh, business, our pharma business, some of these are about 2-3 years old. When we started, we said we will solve for tuberculosis screening, and we said that TB screening is one of those problems that is around. I mean, TB is about 11 million patients every year. 1.5 million patients every year die from TB. It's a completely curable disease. It doesn't get diagnosed on time, and one of the reasons that it doesn't get diagnosed on time is the first line of diagnosis for TB is a chest X-ray. And that chest X-ray that is taken is not read because there is no radiologist available in parts of Africa, parts of India, parts of Southeast Asia. Our first customer, which was a mobile van-based screening program in Philippines, it took them about 4 to 5 weeks to get an X-ray read. And we were able to make that happen— this is in 2018-2019— we were able to make that X-ray read happen in about 30 seconds. So from 4 to 5 weeks to 30 seconds, and that is where value was created. And based on that, WHO came out and said that AI can be used autonomously in the detection of TB on a chest X-ray. And this was guidance that came out in 2021. And since then, I mean, no program in the world is using humans to read X-rays anymore. All X-ray interpretations for TB are done autonomously, uh, with the help of AI, right? And we have done this. It really — you can see, I mean, fields where people are— people go in boats with an X-ray machine and take an AI. AI is sitting on that laptop. You can see, when X-ray is the X-ray machine, you are— this is the laptop where AI is available. You take an X-ray, within 20-30 seconds you have the result. We did that at the Mount Everest Base Camp where we took an X-ray machine. It's like a 20-kilo X-ray machine and the laptop which has AI, and we did screening at that altitude as well. And just to show that it can be done anywhere in the world, and it is now deployed across, I would say, 100 plus countries, 3,000-4,000 plus hospitals, which are hospitals or sites where screening is happening, and finally helping TB get diagnosed earlier and bring down the time of diagnosis from 2 weeks to 2 hours. And in that process, we also save a lot of money because finally it takes— I mean, a total— I mean, there's cost of screening, cost of transporting X-rays, cost of doing the follow-up testing. There's a lot that we're able to save here. And going beyond that, we also added more. I mean, so Gates Foundation funded us to build X-rays, the standard of care for TB screening right now. But Gates Foundation sent a point-of-care ultrasound device. A point-of-care ultrasound is like a, like a razor, and it can fit in your pocket. And we are working on algorithms that can detect TB on a point-of-care ultrasound. Uh, and then we also built a voice, um, uh, a note-taker for doctor-patient conversation. So the idea is that a lot of healthcare workers are solving, uh, problems in these, uh, in these low and middle-income country settings. And how can you upskill those healthcare workers? How can you provide the right nudge to them? And what we have this tool called Aira, which is sitting on a mobile phone, sitting on a tablet, basically listening into the conversation between the healthcare worker and the patient, and it provides the right nudge. So whenever we feel like there is information that can be provided to the healthcare worker, the AI will say, "Okay, this is what you should do with this patient. This is the this is the symptom that you should ask for." Or this is a treatment plan that you should recommend, or this is an investigation that you should recommend. But it provides nudges, it converts that conversation into data, and both are valuable. The nudge itself is super valuable, but even the fact that you are getting data from all these 3,500, 4,000 sites, from millions of healthcare worker conversations, that is very valuable because you don't really know where TB is spreading, which areas are sort of Could be could be pandemic, could be which which diseases could be a pandemic. So again, this kind of tool is being used by several countries as a pandemic preparedness tool. So that's again a significant use case for the for the IRAS machine. And lastly, just wanted to say that all of this comes together in terms of data. I mean, we process about fifteen twenty million patients through our system on an annual basis. So we're getting data scans from those patients, medical records, and then also the actions right. A radiologist sees our sort of recommendation; they might accept it, reject it. So all the clicks, every click that a radiologist does, every click that a healthcare worker does, or physician does, all of that comes back, and there is that flywheel effect of the data that we're getting back from these from these recommendations, and creates a sort of Very sort of the AI model keeps feeding itself. And, uh,

and this is the last slide. Uh, we are—I spoke about 26 FDA clearances. FDA started, but if you look at Europe, we are cleared for 65+ findings. Um, we reach, uh, around—we have reached around 45 million patients in our history, uh, over the last 5 years. Very, very strong scientific, uh, evidence. We published in Nature, Lancet, WHO has recommended our solution for TB screening, and 250+ conference abstracts, evidences, 40 patents, very good publications. And again, a full-stack product which goes from multimodal data, looks at multimodal data from images to text to voice. All of it comes together to help physicians and healthcare workers make the right decisions for the patient. Thank you.

Host (27:28)

Yeah, thanks so much. That was a wonderful and very insightful session. Um, excellent. Lots of questions. I think, uh, people are very curious to find out 2-3 different aspects from the presentation, so I will club it and accordingly ask you the questions. Uh, one question broadly under thematic is more about validations and liabilities and false negatives. So people have asked questions about what is the liability of false negatives for treating lung cancer, what are the kind of EHR liabilities. So maybe we can take that bucket.

Prashant Bhai (28:04)

So the liability is always with the physician that is making that decision, right? The AI is supporting a decision, uh, and even in an— where it's autonomous, finally a healthcare worker organization making the decision, uh, and the liability is almost always on, right? At this point, I would not think that there is any AI, uh, anywhere that takes liability for any software. In fact, not only AI, but any software that takes liability for decisions that the algorithm makes. Uh, I think the only place where we're seeing some of that happen is with autonomous vehicles, and that was pretty limited. So, but that's, that's the only place where AI is लाइटली टेकिंग सम लाइबिलिटी, बट अगेन दैट्स अ दैट्स एन ओपन क्वेश्चन ऑन लॉन्ग टर्म इफ़ एआई इज़ पूरी ऑटोनोमस, कोई गोइंग टेक लाइबिलिटी फॉर दो डिसीज़न्स राइट नाउ राइट नाउ द आंसर फॉर हेल्थकेयर इन जनरल इज़ दैट दी पेशेंट इस फिर रेस्पONSिबल, रेडियोलॉजिस्ट, हेल्थक So as we're building the model, one of the questions is what about the PHI privacy acceptance and the lack of data therefore as we're developing? So PHI typically gets stripped. I mean, all of these decisions can be made without knowing the patient name or, um, I mean, patient information is not relevant to make a decision. If you're looking at a scan, that scan does not require PHI for sort of process it. So most places, uh, we don't use PHI, but PHI, at least the patient identifier, is needed to link back to the EMR record, link back to the fax record. And that is something that we post in the US. Again, there are so many rules around this. Every country has its own data laws, data privacy laws, and we have to comply with that, especially when we operate in so many countries. We have servers which are running within the country. So US gets help from the US AWS cluster, India gets help from India AWS cluster. There are places where there is no AWS, so we can use Alibaba. Malaysia gets help from Malaysia's Alibaba cluster. In the UAE, we work with— so again, I mean, there are different clusters that we— in Vietnam, we work with VSN. So we have to comply with the data privacy laws in every country. And we ensure that we are compliant with those laws. And finally, in the US there is HIPAA, there is GDPR in Europe, every country has their own. And as an AI company, you have to comply. You have to ensure that you have frequent audits for those and ensure that the audit— you are basically compliant to the laws.

Host (30:50)

And just one final question on this bucket. Regulatory risk for the growth of your company, are there any examples specifically? So regulatory risk, as you see, is nothing more than—

Prashant Bhai (31:04)

I think regulatory risk is little. I mean, because we in many ways are— our model is to work with physicians and provide a diligent support pool, medical diligent support pool to physicians. Where you are going outside of physician control and you're saying that I'll provide decision directly to the patient, those are areas where there is a related risk in my mind associated with AI. But if you work with physicians, I don't think there is much risk right now.

Host (31:33)

Uh, we'll switch gears and move a little bit towards economics now. Take one last question. Uh, you showed your presence from one of the best hospitals. What's the deployment get into? That's the first question.

Prashant Bhai (31:46)

See, the, um, the price points of deployment in India are probably about 10 to 20x lower than what we get paid in

the US, right? And we deploy, we work with a lot of health systems in India, but again, our go-to market is not in India. Go-to market is in the US, and go-to market in India is only around tuberculosis. It is not for cancer. Even the technology adoption, even simple things like EMR, right? A lot of our tool requires EMR data to take the action out, to sort of make inferences, to action out the follow-ups, to provide guideline-concordant follow-ups. The value that we create is mostly for systems like the US where there is EMR, broad EMR adoption, where there are guideline-concordant follow-ups. I think in India the value is lower for these kinds of screenings. TB screening is still very valuable, but outside of that, I think the value is lower also.

Host (32:44)

And in terms of value creation, therefore, uh, what has been the maximum driver of this value creation? Insurance companies, hospital themselves? How does that economic code—

Prashant Bhai (32:55)

So the value, again, I mean, this is, uh, very interesting. I also learned as I built up here, right? The value of a single lung cancer patient to a hospital is about 100K. If I recommend that this patient needs to be followed up, there is 100K worth of services that the hospital can provide that patient, from multiple CT scans to a biopsy to surgery to sort of other patient doctor consultations, physician consultations in between, for the 100K worth of value. The value of a lung cancer patient to a pharma company is \$200,000 plus because one month of drugs for a lung cancer patient is about \$17,500. This is across pharma. So if I look at AstraZeneca's drugs, GMS, Johnson Johnson, and so if you think about a patient taking therapy for 1 to 2 years, it's easily \$200,000 to \$400,000. So there is tremendous value created when you diagnose early because patients are on therapy longer. Patients can get surgeries, for example. So the sort of economics of that early diagnosis is not just for the patient; it is also for the pharma companies and the hospitals. Interesting.

Host (34:09)

And now that people who are resource fee are democratizing clinical interventions at great length, is that a threat for this monetization of your product? Not really.

Prashant Bhai (34:21)

I mean, first of all. First of all, the open models are very general purpose, and if you look at any sort of frontier model, any open source model, they will not do difficult tasks such as finding a small nodule behind the heart or behind a rib easily. They will not do that task well, and we are— we have millions of scans of that particular kind which help us do very well on that specific task. So I think the general purpose models will well on general stuff, but to solve really niche problems, difficult problems in healthcare, they're not that good, number one. But also, second thing is that, um, all of these require FDA approval. So if you use them in a system, you have to get an FDA approval. Again, they're not approved for general use, that's second thing. And the third thing, which is probably most important, is the fact that none of this works if you don't integrate well with the hospital. That integration is a huge challenge. I mean, EMR integration impacts innovation. It is a huge amount of work, and that is something where it's not like— I mean, people will just start, okay, I'll just start using this, this model, it's really important, I'll start deploying it in the hospital. Like you said, there are so many data privacy settings and so many things, right, where it's not easy for physicians to use a technology. They cannot just take a scan and send it to ChatGPT. So there are, there are, uh, challenges there. But I think across— I mean, the board, I would say that especially for niche problems, AI is going to be better than— sorry, the special purpose AI is going to be better than general models.

Host (35:50)

You spoke about a very important element, integration, which has also to do with your cost line item. Can you explain a little bit about how this integration works in the US? Who does it for you? And any difference versus other geographies?

Prashant Bhai (36:01)

So we do it, we do it, but the integration is mainly about— first of all, I mean, they have quite a lot of data compliance regulations, they have architecture reviews, um, there is the prior integration with EMR data, there is integration with PACS. So all of those require a lot of effort both from the hospital IT as well as our team, sometimes from the vendor that is providing that as well. So for example, Epic is the EHR system, or whatever is the EHR system, from them as well. So getting in and deployed in the system. And getting access to our data and

being able to provide access back, bidirectional integration, is a massive challenge. And it takes time. It takes 3 to 4 months to get all of that to happen. It's not actual like hours of work. Hours of work might be like a week or 2 hours, for example. But just getting time is ensuring that these are done in a certain fashion.

Host (36:58)

सीकरेटेड वर्ल्ड लॉन्च | मैंने इंटरैस्टिंग कमेंट करा हुआ है, पॉफिट | सो इफ यू हैव गिवन वन डॉलर ऑफ इंफ्रीमेंटेड कैपिटल टू यू, व्हेन वुड यू डिप्लॉय इट? क्वेश्चन सेम मिली तू जेनरेट बेटर लाइफटाइम वैल्यू, मोर अल्लोरिदम्स, मोर डिस्ट्रीब्यूशन, लेनिकल वर्कफ़्लो? मोर That's very clear. In fact, then one more lens. I think intuitive program with circular. Can you think?

Prashant Bhai (37:39)

So intuitive surgical. I don't know. Maybe that question is actually mixed.

Host (37:44)

Maybe that's me. Let's separate it out. There is intuitive surgical and then there is circular.

Prashant Bhai (37:48)

So with circular, we are working on early COPD. Detection across Indian hospital system, and we deploy with them to diagnose COPD on incidental patients. X-rays get taken routinely for so many reasons. On those routine X-rays, we are finding risk of COPD and triaging those patients for follow-up. So that's happening with Sigla. With Intuitive, the partnership is that they have a very strong distribution. We also distribute to pulmonologists because Those are the physicians that we sell to, and Intuitive also distributes to pulmonologists. They have this device called the Ion device, which does a biopsy of bronchoscopy. And what we sort of agreed on about a year ago is that we'll benefit from them distributing our products to the pulmonologists. They have very, very strong relationships with US pulmonology departments, and they then distribute our software along with their hardware products. So both are being distributed together. And finally, I mean, the deployment, integration, everything comes to us, but they sort of build the initial lead. They speak about our product with the terminology team, and then we get involved, our sales team gets involved, and then we sort of finish the transaction, finish the contracting, and then deploy. So that's the partnership with Intuitive in the US.

Host (39:06)

That's perfect. Application. Uh, do you plan to build beyond diagnostics AI into, say, molecular diagnostics, biomarkers?

Prashant Bhai (39:18)

So we are looking at, uh, like I said, 3 areas. One is multimodal AI, and multimodal includes everything. But right now, the idea of multimodal is that it is looking at, uh, voice, text, and image. It's not an image, is mostly radiology image right now, but the text itself includes a lot of pathology data. So you could look at pathology reports. Pathology reports might say that this patient is EGFR positive, for example, one of the one genetic marker, right? So you could pull that data. You don't have to process the pathology images to be able to pull that data. You can just look at the pathology report. Okay. So but we do we do ingest a lot of different data sources. We have full patient context, I mean from everything, full patient history, from all the different tests they have done and the results from that, along with the radiology images which we can process and identify things that the radiologist might have missed. And based on all of that, we are then able to recommend that this patient needs a follow-up. Then the second piece is ensuring that the follow-up happens, the follow-up management too. And third is the area where we are actually ग्रोइंग वेरी रैपिडली इस तू प्रोवाइड द राइट वैल्यू एंड कंपेयरेंट ट्रीटमेंट ऑफ़ डिसेम्बलेशन एंड दैट अगेन इस सिग्निफिकेंट वैल्यू फॉर द सिस्टम और फॉर फार्मा कंपनीज़ | आई थिंक इन द पॉसिबल टाइम आई विल कट लास्ट क्वेश्चन, देर आर सो मेनी अदर्स में जस्ट गोज़ ओवर द कॉफ़ So training doctors is hard, I mean, actually, because they don't have a lot of time, right? So you have to make the product super easy to use. You have to have pop-ups that come at the right time. For example, radiologists, when they use it, they will read normally, and when there is something that Cure has pointed out, it will show us a pop-up on the screen that this is a nodule, there is this patient that has a लेटली सो पी और अब भी जैसे पॉपअप, देखें क्लिक ऑन दैट, इट ओपन अप अ स्मॉल अ बिगर पॉपअप विच शोज़ एक्जेक्टली वेर दैट अपार्मरिटीज़ में डिटेल्स मोर अबाउट दैट, एंड देन देखें इफ़ दे वांट, देखें गोइंग तू अ होल स्क्रीन वेर द प्रोवाइड दे कुछ कैन रिव्यू एवरीथिंग, बट वी म सीक्रेट सीक्रेट वर्क | तो दैट्स व्हाट आई टॉकड अबाउट इन इंटीग्रेशन, दैट यू कैन नॉट क्रिएट मोर वर्क फॉर एनीबडी | इफ़ यू क्रिएट मोर वर्क, नोटिस में यूज़ इट | सो वी ट्राई टू मिनिमाइज़ द अमाउंट ऑफ़ वर्क एनीबडी हैज़ टू डू तू यूज़ अस परिशन | एवरीथिंग इज़ इंटीग्र They can open it up, see what is available on our page, and they can look at the votes. And on, on the end, we have a full, uh, what is called a DICOM viewer. So you can see an

image, the full DICOM viewer, they can see all of the abnormalities on that. But typically we don't— training is simpler, I can give done in a day, and we don't add a lot of load to the application.

Host (42:30)

I think thank you so much for that insightful session and answering all the